# Graded word similarity in context by composing static and contextual embeddings

Tim Lawson

December 16, 2023

## 1    Introduction

In his *Foundations of Arithmetic*, Frege promises "never to ask for the meaning of a word in isolation, but only in the context of a proposition" (1980, p. xvii). This 'context principle' is intuitive: words are frequently polysemous, or assume different connotations and emphasis within different expressions. Historically, however, contextuality has been a problem for distributional meaning representations. Founded on the distributional hypothesis (Harris 1954; Firth 1957), both count-based and predictive models of word meaning[1] originally produced a single representation for each word in the model's vocabulary. One of these *static* representations must, therefore, encode all of a word's senses and connotations, which may obstruct its use in modelling context-dependent phenomena.

Prior to the widespread availability of pre-trained word embeddings (e.g., Mikolov, Chen, et al. 2013; Pennington et al. 2014) and their successors, this problem was generally addressed by one of two approaches: firstly, by producing a representation for each sense of a target word and disambiguating between them in the given context (*word-sense disambiguation*); or secondly, by composing the representation of the target word with the representations of the words in its context (*contextualisation*). These approaches have been largely overshadowed by the advent of model architectures that take sequences as inputs and naturally produce *contextual* representations of the items in the sequence, such as Transformers (Vaswani et al. 2017).

To my knowledge, however, there has been scant direct comparison of the performance of these contextual representations with the application of prior methods of contextualisation to static representations. SemEval-2020 Task 3, "Graded Word Similarity in Context" (Armendariz, Purver, Pollak, et al. 2020), presents an opportunity to make such a comparison. Briefly, the task is to predict the human judgment of similarity of the same pair of words in two different contexts (section 2). I elected to focus on the first sub-task, which is to predict the *change* in similarity, rather than the absolute similarity in each context. Specifically, I evaluated the results obtained by computing the cosine similarity between static and contextual embeddings and the composition of these embeddings within a fixed-size context window.

## 2    Task definition

The first sub-task of SemEval-2020 Task 3 is to predict the direction and magnitude of the change in the human judgment of similarity of the same pair of target words in two different contexts. The task is unsupervised: the submissions were evaluated on the CoSimLex dataset (Armendariz, Purver, Ulčar, et al. 2020, pp. 39–42) but only a minimal 'practice kit' of fewer than ten instances was provided in advance. CoSimLex is an extension of SimLex-999 (Hill et al. 2015) that consists of pairs of target words and their contexts in four languages: English ($n = 340$), Finnish ($n = 24$), Croatian ($n = 112$), and Slovene ($n = 111$).

---

[1]This terminological distinction is due to Baroni et al. (2014).

The score for the first sub-task was computed by the uncentered (zero-mean) Pearson correlation coefficient between the predicted changes in similarity and the human judgments represented in the CoSimLex dataset (Armendariz, Purver, Ulčar, et al. 2020, p. 42). This metric is equivalent to the cosine similarity between the two vectors of results:

$$\text{score}(\vec{\hat{y}}, \vec{y}) = \frac{\sum_{i=1}^{n} \hat{y}_i y_i}{\left(\sum_{i=1}^{n} \hat{y}_i^2\right)\left(\sum_{i=1}^{n} y_i^2\right)} = \frac{\vec{\hat{y}} \cdot \vec{y}}{\|\vec{\hat{y}}\|\|\vec{y}\|} \tag{1}$$

A consequence of this choice of evaluation metric is that composing representations by addition or the arithmetic mean produces the same results; hence, I elected to only evaluate addition between them.

## 3   Related work

Many context-based approaches to word-sense disambiguation have been proposed since the advent of count-based models of word meaning. In Landauer and Dumais (1997)'s introduction of Latent Semantic Analysis (LSA), the authors argued that taking the average of the high-dimensional representation of a word with those of its immediate context may be sufficient to determine the word's contextual meaning (ibid., pp. 229–230). The contextualisation of representations of word meanings is intimately related to their composition to form representations of phrase and sentence meanings. This connection is evident in Kintsch (2001)'s proposal of a procedure to modify the vector of a predicate according to its argument in a given context. Mitchell and Lapata (2008) also adapted Kintsch's demonstration to evaluate alternative composition operations.

Vector addition and averaging continue to be 'surprisingly effective' means to compose word embeddings (Boleda 2020, p. 10), and addition produces plausible results for the word–analogy task[2] (e.g., Mikolov, Chen, et al. 2013, p. 9; Mikolov, Sutskever, et al. 2013, p. 7).

Nevertheless, models that produce contextual embeddings have achieved widespread success on benchmarks that involve language understanding. As Arora et al. (2020) point out, this comes at a significant computational cost, and static embeddings may be sufficient for many tasks. Furthermore, Gupta et al. (2019) and Bommasani et al. (2020) have shown that static embeddings can be obtained from contextual-embedding models, which outperform the embeddings of static models while retaining their computational advantages.

> 1: Finish this section.

## 4   Methodology

Originally, it would not have been possible to optimise a parameterised model for the task except by reference to a separate dataset; therefore, I chose to focus on the application of pre-trained static and contextual embeddings. The basic procedure of the methods that I evaluated is as follows. For each pair of target words and each of the two contexts in which they appear, I obtained a contextualised representation of a target word by: finding the index of the target word's first sub-word token within the tokens of the target word's context; finding the tokens within a fixed-size window around the target word's first token; obtaining the embeddings of the tokens in the window; and combining the the embeddings to produce a single representation of the target word.

My primary motivation in undertaking this task was to compare static and contextual embeddings. However, several of the submissions to SemEval–2020 Task 3 used a combination of the weights of a Transformer model's hidden states (e.g., Gamallo 2020, p. 276; Pessutto et al. 2020, p. 3; Hettiarachchi and Ranasinghe 2021, p. 4). Therefore, I chose to also evaluate an example of this approach. A thorough comparison of the performance of its variants is beyond the scope of this paper. In all cases, the tokenization was performed by and the embeddings were obtained from pre-trained models

---

[2]Also known as the parallelogram model of analogy (Rumelhart and Abrahamson 1973).

…from the Isenj. He <u>accepts</u> their offer but knows that …

$\downarrow$    1. Find first sub-word token

$$\left[ \dots, \text{\#\#nj}, ., \text{he}, \underline{\text{accept}}, \text{\#\#s}, \text{their}, \text{offer}, \dots \right]$$

$\downarrow$    2. Find tokens in a fixed-size window

$$\left[ \text{he}, \text{accept}, \text{\#\#s} \right]$$

$\downarrow$    3. Obtain representations

$$\left[ \overrightarrow{\text{he}}, \overrightarrow{\text{accept}}, \overrightarrow{\text{\#\#s}} \right]$$

$\downarrow$    4. Combine representations

$$\overrightarrow{accept} = \overrightarrow{\text{he}} + \overrightarrow{\text{accept}} + \overrightarrow{\text{\#\#s}}$$
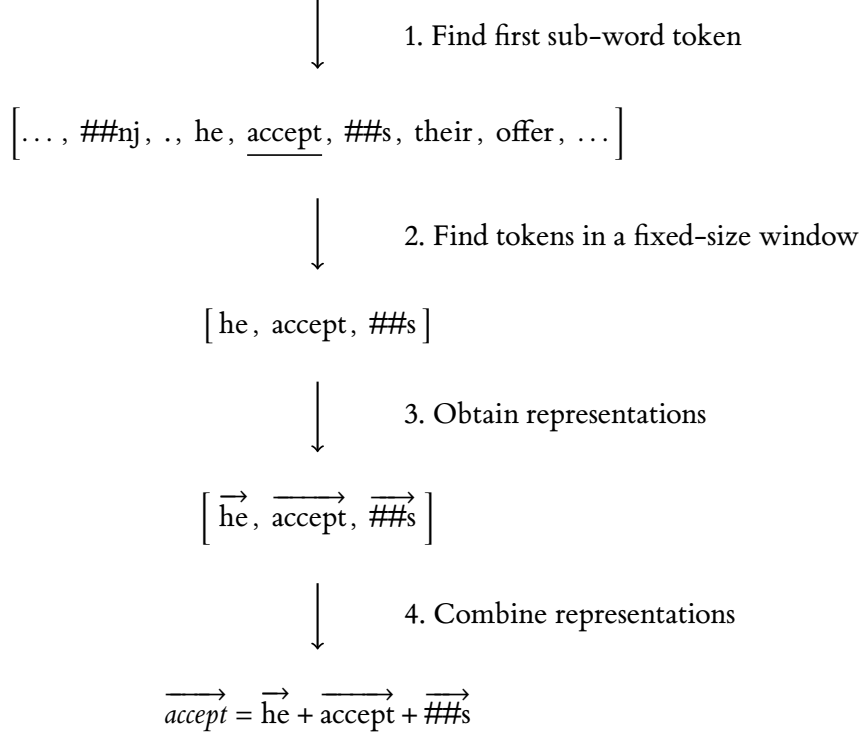
Figure 1: A schematic of the procedure used to obtain a contextualised representation of a target word from pre-trained static or contextual embeddings. In this example, the target word is "accept", the window size is one (either side of the target word), and the composition operation is addition.

| Model name | English | Finnish | Croatian | Slovene |
|---|---|---|---|---|
| EMBEDDIA/crosloengual-bert[1] | ✓ | ✓ | ✓ | ✓ |
| TurkuNLP/bert-base-finnish-cased-v1[2] | | ✓ | | |
| TurkuNLP/bert-base-finnish-uncased-v1[2] | | ✓ | | |
| TurkuNLP/bert-large-finnish-cased-v1[2] | | ✓ | | |
| bert-base-cased | ✓ | | | |
| bert-base-multilingual-cased | ✓ | ✓ | ✓ | ✓ |
| bert-base-multilingual-uncased | ✓ | ✓ | ✓ | ✓ |
| bert-base-uncased | ✓ | | | |
| bert-large-cased | ✓ | | | |
| bert-large-cased-whole-word-masking | ✓ | | | |
| bert-large-uncased | ✓ | | | |
| bert-large-uncased-whole-word-masking | ✓ | | | |
| classla-bcms-bertic[3] | | | ✓ | |

Figure 2: The pre-trained models available via the HuggingFace *Transformers* library (Wolf et al. 2020) that I chose to evaluate for each language. The corresponding references are [1]Ulčar and Robnik–Šikonja (2020), [2]Virtanen et al. (2019), [3]Ljubešić and Lauc (2021), and Devlin et al. (2019) otherwise.
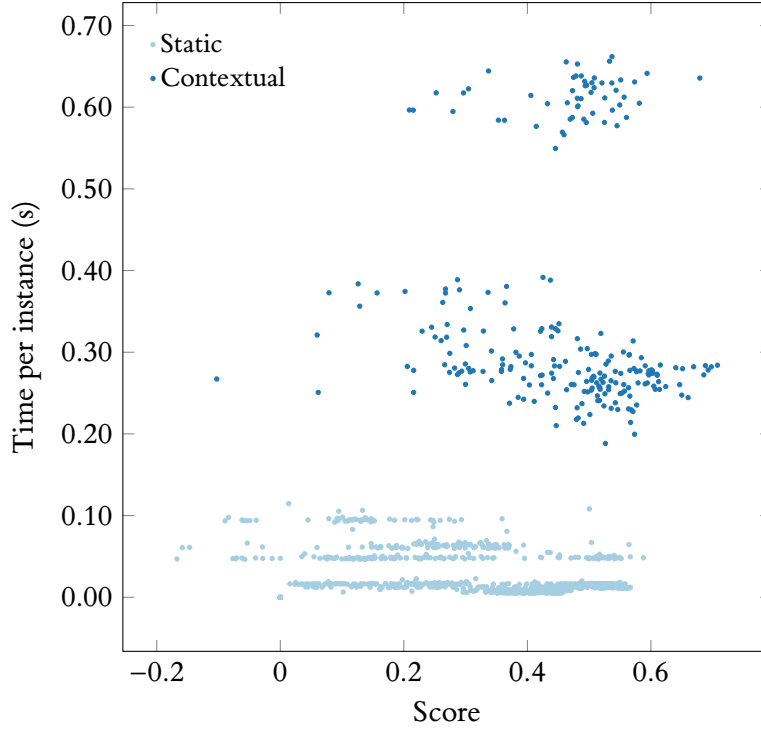
Figure 3: The time per instance against the score obtained by different models and languages with the additive composition operation. The static-embedding models take less time per instance and show less variability. The `large` variants of the contextual-embedding models take the most time.

available via the HuggingFace *Transformers* library (Wolf et al. 2020). The models that I evaluated for each language are given in fig. 2. For each of these, I investigated the following variants:

- `static`: the model's input embeddings

- `contextual`: the model's output embeddings

- `pooled`: the sum of the model's last four hidden-states

Notably, the use of a sub-word vocabulary by these models (e.g., Devlin et al. 2019, p. 4174) dictates that a target word may be represented by a different number of tokens in each context. As a result, the representations of a pair of target words may be different in each context, even if they are static and the window size is zero. This is the cause of the non-zero scores obtained by models of this kind, particularly for the Finnish language.

Inspired by Landauer and Dumais (1997), Kintsch (2001), and Mitchell and Lapata (2008), I predominantly investigated the application of element-wise addition and multiplication as composition operations. However, preliminary experiments indicated that multiplication performed poorly across all languages, models, and window sizes; hence, it was discarded before the final analysis. Additionally, I chose to evaluate the concatenation ('stacking') of embeddings. In the case that the number of embeddings was fewer than the context-window size, i.e., the target word was close to the beginning or the end of its context, I right-padded the concatenated embedding with zeros to obtain contextual embeddings of equal length.

## 5 Results

### 5.1 Cost-benefit analysis of contextual embeddings

2: Write this section.

## 5.2 Language-specificity of window-size effects

Generally, I found that the scores obtained by both static- and contextual-embedding models were maximised by a non-zero context-window size. Due to the computational expense of exhaustively searching the possible window sizes, I applied a heuristic to constrain the search space. A naïve estimation of the average number of words in each context, i.e., segmenting on whitespace, gave a result of between 40 and 60 for the different languages. Therefore, for the static-embedding models, I chose 50 as an upper bound on the window size on either side of the target word. The motivation to choose a smaller maximum window size for contextual-embedding models was similarly economical (section 5.1); however, as the window size approaches the length of the sequence, one would expect a combination of token representations to be superseded by the sequence-level representation of the model, e.g., the special CLS token of BERT variants (Devlin et al. 2019, p. 4174). These heuristics were largely vindicated by the results of the evaluation, which demonstrated that the scores decrease as the window size approaches the maximum.

The influence of the window size is intuitive in the case of static embeddings. Without a context window, the representations of a target word only differ between expressions if the word is represented by different sub-word tokens in the different expressions. A similar argument applies to contextual embeddings, in that a target word may be represented by multiple sub-word tokens. The window sizes that maximise the score for each language and model are given in figs. 4 and 5.

## 6 Discussion

Batchkarov et al. (2016) critically analyse word similarity as an evaluation methodology for distributional semantic models. In particular, the notion of 'similarity' manifested by these models encompasses a broad range of semantic relations (e.g., Padó and Lapata 2003, p. 2), with the consequence that performance on an intrinsic word-similarity task does not necessarily translate to extrinsic downstream tasks (Batchkarov et al. 2016, pp. 7–8). Moreover, inter-annotator agreement is generally poor for word-similarity tasks in comparison to more specific downstream tasks (ibid., pp. 8–9).

> 3: Discuss results in the above context.

## 7 Conclusion

In this paper, I have presented the results of a hypothetical submission to SemEval-2020 Task 3, "Graded Word Similarity in Context". The purpose of this evaluation was to compare the performance of static and contextual embeddings and their composition within a fixed-size context window on the task of predicting the change in the human judgment of similarity of a pair of words in two different contexts. I found that contextual embeddings generally outperformed static embeddings, but at a significant computational cost, and that composition benefited both static and contextual embeddings of sub-word tokens, with highly language-specific dependence on the window size.

The results that I have given must be interpreted in context: the original submission authors did not have access to the evaluation dataset prior to submitting their results, only the 'practice kit' of very few instances, and were therefore unable to optimise a parameter such as the window size prior to submission. Additionally, the models that I used were not necessarily available to the authors. With these caveats, I achieved several notable results:

- The pooled variant of EMBEDDIA/crosloengual-bert with a window size of three would have placed second among the Croatian submissions, with a score of 0.717.

- The contextual variant of TurkuNLP/bert-base-finnish-uncased-v1 with a window size of zero would have placed fourth among the Finnish submissions, with a score of 0.679.

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 17 | 0.439 |
| bert-base-cased | 18 | 0.456 |
| bert-base-multilingual-cased | 18 | 0.455 |
| bert-base-multilingual-uncased | 19 | 0.455 |
| bert-base-uncased | 14 | 0.442 |
| bert-large-cased | 18 | 0.441 |
| bert-large-cased-whole-word-masking | 18 | 0.465 |
| bert-large-uncased | 16 | 0.440 |
| bert-large-uncased-whole-word-masking | 16 | 0.471 |

(a) English

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 21 | 0.588 |
| TurkuNLP/bert-base-finnish-cased-v1 | 0 | 0.452 |
| TurkuNLP/bert-base-finnish-uncased-v1 | 0 | 0.564 |
| TurkuNLP/bert-large-finnish-cased-v1 | 0 | 0.500 |
| bert-base-multilingual-cased | 3 | 0.394 |
| bert-base-multilingual-uncased | 29 | 0.369 |

(b) Finnish

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 31 | 0.550 |
| bert-base-multilingual-cased | 32 | 0.535 |
| bert-base-multilingual-uncased | 31 | 0.558 |
| classla/bcms-bertic | 31 | 0.567 |

(c) Croatian

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 11 | 0.383 |
| bert-base-multilingual-cased | 8 | 0.316 |
| bert-base-multilingual-uncased | 10 | 0.291 |

(d) Slovene

Figure 4: The window size that maximises the score for the `static` model variants with the additive composition operation. The best score for each language is underlined.

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 3 | 0.607 |
| bert-base-cased | 1 | 0.556 |
| bert-base-multilingual-cased | 1 | 0.574 |
| bert-base-multilingual-uncased | 3 | 0.577 |
| bert-base-uncased | 1 | 0.660 |
| bert-large-cased | 3 | 0.538 |
| bert-large-cased-whole-word-masking | 2 | 0.557 |
| bert-large-uncased | 1 | 0.521 |
| bert-large-uncased-whole-word-masking | 1 | 0.594 |

(a) English

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 10 | 0.366 |
| TurkuNLP/bert-base-finnish-cased-v1 | 1 | 0.615 |
| TurkuNLP/bert-base-finnish-uncased-v1 | 1 | 0.483 |
| TurkuNLP/bert-large-finnish-cased-v1 | 1 | 0.679 |
| bert-base-multilingual-cased | 3 | 0.519 |
| bert-base-multilingual-uncased | 2 | 0.423 |

(b) Finnish

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 3 | 0.708 |
| bert-base-multilingual-cased | 5 | 0.614 |
| bert-base-multilingual-uncased | 6 | 0.603 |
| classla/bcms-bertic | 10 | 0.423 |

(c) Croatian

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 3 | 0.542 |
| bert-base-multilingual-cased | 3 | 0.572 |
| bert-base-multilingual-uncased | 9 | 0.439 |

(d) Slovene

Figure 5: The window size that maximises the score for the contextual model variants with the additive composition operation. The best score for each language is underlined.

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 1 | 0.617 |
| bert-base-cased | 1 | 0.629 |
| bert-base-multilingual-cased | 1 | 0.565 |
| bert-base-multilingual-uncased | 1 | 0.629 |
| bert-base-uncased | 1 | <u>0.653</u> |
| bert-large-cased | 1 | 0.624 |
| bert-large-cased-whole-word-masking | 3 | 0.568 |
| bert-large-uncased | 3 | 0.624 |
| bert-large-uncased-whole-word-masking | 1 | 0.574 |

(a) English

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 10 | 0.343 |
| TurkuNLP/bert-base-finnish-cased-v1 | 1 | 0.632 |
| TurkuNLP/bert-base-finnish-uncased-v1 | 1 | 0.562 |
| TurkuNLP/bert-large-finnish-cased-v1 | 1 | <u>0.668</u> |
| bert-base-multilingual-cased | 2 | 0.466 |
| bert-base-multilingual-uncased | 2 | 0.586 |

(b) Finnish

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 3 | <u>0.717</u> |
| bert-base-multilingual-cased | 6 | 0.583 |
| bert-base-multilingual-uncased | 1 | 0.631 |
| classla/bcms-bertic | 2 | 0.662 |

(c) Croatian

| Model name | Window size | Score |
|---|---|---|
| EMBEDDIA/crosloengual-bert | 2 | <u>0.544</u> |
| bert-base-multilingual-cased | 3 | 0.453 |
| bert-base-multilingual-uncased | 2 | 0.412 |

(d) Slovene

Figure 6: The window size that maximises the score for the pooled model variants with the additive composition operation. The best score for each language is underlined.

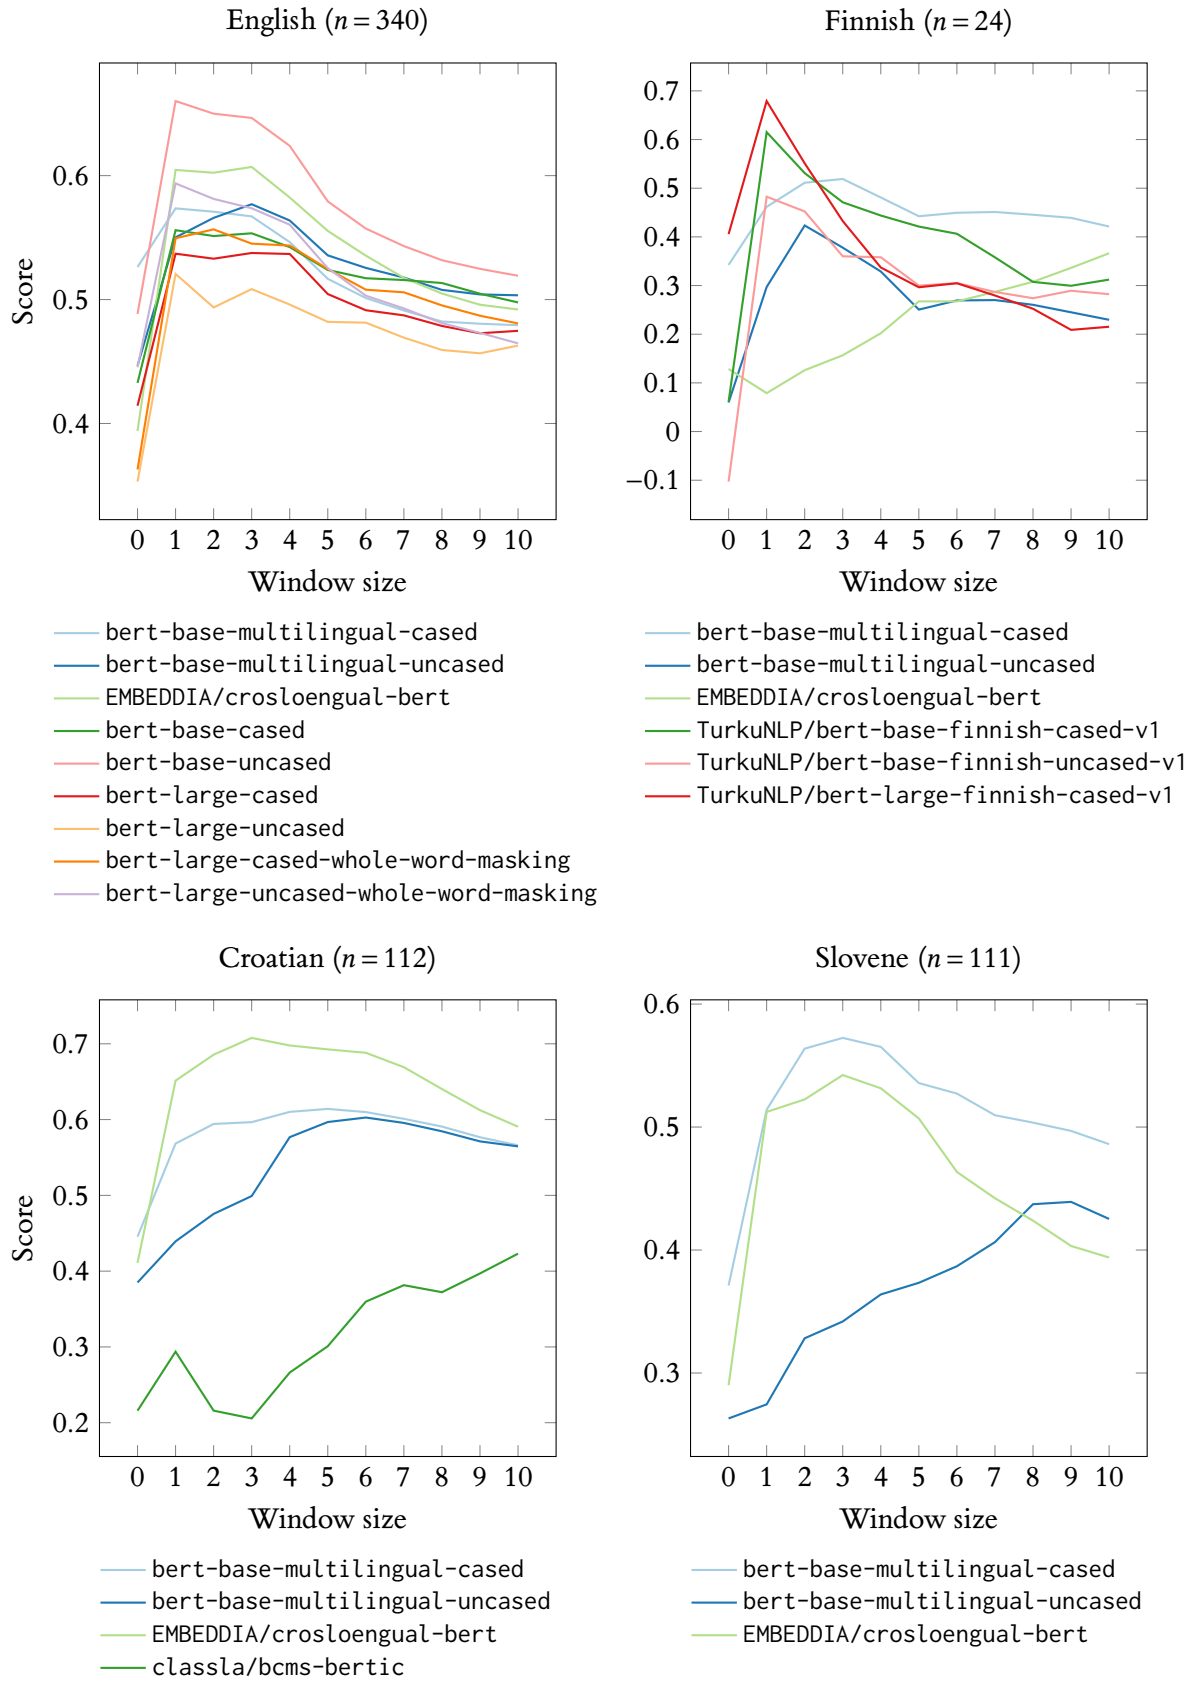Figure 7: The score against window size for the static model variants with the additive composition operation.

Figure 8: The score against window size for the contextual model variants with the additive composition operation.

English ($n = 340$)

Finnish ($n = 24$)

bert-base-multilingual-cased
bert-base-multilingual-uncased
EMBEDDIA/crosloengual-bert
bert-base-cased
bert-base-uncased
bert-large-cased
bert-large-uncased
bert-large-cased-whole-word-masking
bert-large-uncased-whole-word-masking

bert-base-multilingual-cased
bert-base-multilingual-uncased
EMBEDDIA/crosloengual-bert
TurkuNLP/bert-base-finnish-cased-v1
TurkuNLP/bert-base-finnish-uncased-v1
TurkuNLP/bert-large-finnish-cased-v1

Croatian ($n = 112$)

Slovene ($n = 111$)

bert-base-multilingual-cased
bert-base-multilingual-uncased
EMBEDDIA/crosloengual-bert
classla/bcms-bertic

bert-base-multilingual-cased
bert-base-multilingual-uncased
EMBEDDIA/crosloengual-bert

Figure 9: The score against window size for the `pooled` model variants with the additive composition operation.

- The `static` variant of `TurkuNLP/bert-base-finnish-uncased-v1` with a window size of zero outperform several of the Finnish submissions, including the baseline, with a score of 0.564.

Given the significant expense of applying contextual-embedding models, these results highlight the importance of analysing the complexity of the task at hand and considering the possibility that a simpler model produces adequate results.

# References

Armendariz, Carlos Santos, Matthew Purver, Senja Pollak, et al. (2020). "SemEval-2020 Task 3: Graded Word Similarity in Context". In: *Proceedings of the Fourteenth Workshop on Semantic Evaluation*. Barcelona (online): International Committee for Computational Linguistics, pp. 36–49.

Armendariz, Carlos Santos, Matthew Purver, Matej Ulčar, et al. (2020). *CoSimLex: A Resource for Evaluating Graded Word Similarity in Context*.

Arora, Simran et al. (2020). "Contextual Embeddings: When Are They Worth It?" In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Ed. by Dan Jurafsky et al. Online: Association for Computational Linguistics, pp. 2650–2663.

Baroni, Marco, Georgiana Dinu, and Germán Kruszewski (2014). "Don't Count, Predict! A Systematic Comparison of Context-Counting vs. Context-Predicting Semantic Vectors". In: *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Baltimore, MD: Association for Computational Linguistics, pp. 238–247.

Batchkarov, Miroslav et al. (2016). "A Critique of Word Similarity as a Method for Evaluating Distributional Semantic Models". In: *Proceedings of the 1st Workshop on Evaluating Vector-Space Representations for NLP*. Berlin, Germany: Association for Computational Linguistics, pp. 7–12.

Boleda, Gemma (2020). "Distributional Semantics and Linguistic Theory". In: *Annual Review of Linguistics* 6.1, pp. 213–234.

Bommasani, Rishi, Kelly Davis, and Claire Cardie (2020). "Interpreting Pretrained Contextualized Representations via Reductions to Static Embeddings". In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Ed. by Dan Jurafsky et al. Online: Association for Computational Linguistics, pp. 4758–4781.

Devlin, Jacob et al. (2019). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding". In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Ed. by Jill Burstein, Christy Doran, and Thamar Solorio. Minneapolis, MN: Association for Computational Linguistics, pp. 4171–4186.

Firth, J. R. (1957). "A Synopsis of Linguistic Theory, 1930–1955". In: *Studies in Linguistic Analysis*. Oxford, UK: Basil Blackwell, pp. 1–32.

Frege, Gottlob (1980). *The Foundations of Arithmetic: A Logico-Mathematical Enquiry Into the Concept of Number*. 2nd rev. ed. Evanston, Ill: Northwestern University Press.

Gamallo, Pablo (2020). "CitiusNLP at SemEval-2020 Task 3: Comparing Two Approaches for Word Vector Contextualization". In: *Proceedings of the Fourteenth Workshop on Semantic Evaluation*. Ed. by Aurelie Herbelot et al. Barcelona (online): International Committee for Computational Linguistics, pp. 275–280.

Gupta, Prakhar, Matteo Pagliardini, and Martin Jaggi (2019). *Better Word Embeddings by Disentangling Contextual N-Gram Information*.

Harris, Zellig S. (1954). "Distributional Structure". In: *WORD* 10.2-3, pp. 146–162.

Hettiarachchi, Hansi and Tharindu Ranasinghe (2021). *BRUMS at SemEval-2020 Task 3: Contextualised Embeddings for Predicting the (Graded) Effect of Context in Word Similarity*. URL: http://arxiv.org/abs/2010.06269.

Hill, Felix, Roi Reichart, and Anna Korhonen (2015). "SimLex-999: Evaluating Semantic Models With (Genuine) Similarity Estimation". In: *Computational Linguistics* 41.4, pp. 665–695.

Kintsch, Walter (2001). "Predication". In: *Cognitive Science* 25.2, pp. 173–202.

Landauer, Thomas K. and Susan T. Dumais (1997). "A Solution to Plato's Problem: The Latent Semantic Analysis Theory of Acquisition, Induction, and Representation of Knowledge". In: *Psychological Review* 104.2, pp. 211–240.

Ljubešić, Nikola and Davor Lauc (2021). "BERTić – The Transformer Language Model for Bosnian, Croatian, Montenegrin and Serbian". In: *Proceedings of the 8th Workshop on Balto-Slavic Natural Language Processing*. Ed. by Bogdan Babych et al. Kiyv, Ukraine: Association for Computational Linguistics, pp. 37–42.

Mikolov, Tomáš, Kai Chen, et al. (2013). *Efficient Estimation of Word Representations in Vector Space.*

Mikolov, Tomáš, Ilya Sutskever, et al. (2013). "Distributed Representations of Words and Phrases and Their Compositionality". In: *Advances in Neural Information Processing Systems*. Ed. by C. J. Burges et al. Vol. 26. Curran Associates, Inc.

Mitchell, Jeff and Mirella Lapata (2008). "Vector-Based Models of Semantic Composition". In: *Proceedings of ACL-08: HLT*. Columbus, Ohio: Association for Computational Linguistics, pp. 236–244.

Padó, Sebastian and Mirella Lapata (2003). "Constructing Semantic Space Models from Parsed Corpora". In: *Proceedings of the 41st Annual Meeting of the Association for Computational Linguistics*. Sapporo, Japan: Association for Computational Linguistics, pp. 128–135.

Pennington, Jeffrey, Richard Socher, and Christopher D. Manning (2014). "GloVe: Global Vectors for Word Representation". In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Ed. by Alessandro Moschitti, Bo Pang, and Walter Daelemans. Doha, Qatar: Association for Computational Linguistics, pp. 1532–1543.

Pessutto, Lucas R. C. et al. (2020). *BabelEnconding at SemEval-2020 Task 3: Contextual Similarity as a Combination of Multilingualism and Language Models*. URL: http://arxiv.org/abs/2008.08439.

Rumelhart, David E. and Adele A. Abrahamson (1973). "A Model for Analogical Reasoning". In: *Cognitive Psychology* 5.1, pp. 1–28.

Ulčar, Matej and Marko Robnik-Šikonja (2020). "FinEst BERT and CroSloEngual BERT". In: *Text, Speech, and Dialogue*. Ed. by Petr Sojka et al. Lecture Notes in Computer Science. Cham: Springer International Publishing, pp. 104–111.

Vaswani, Ashish et al. (2017). "Attention Is All You Need". In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc.

Virtanen, Antti et al. (2019). *Multilingual Is Not Enough: BERT for Finnish.*

Wolf, Thomas et al. (2020). *HuggingFace's Transformers: State-of-the-art Natural Language Processing.*